

## Solution Assignment 3

wordmean: A map/reduce program that counts the average length of the words in the input files.

```
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordmean /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs/out/wordmeanfile1
18/01/23 19:16:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/01/23 19:16:43 INFO client.RMProxy: Connecting to ResourceManager at localhost:127.0.0.1:8032
18/01/23 19:16:45 INFO input.FileInputFormat: Total input files to process : 1
18/01/23 19:16:45 INFO mapreduce.JobSubmitter: number of splits:1
18/01/23 19:16:45 INFO Configuration.deprecation: yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
18/01/23 19:16:45 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1516707793508_0012
18/01/23 19:16:46 INFO impl.YarnClientImpl: Submitted application application_1516707793508_0012
18/01/23 19:16:46 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1516707793508_0012/
18/01/23 19:16:46 INFO mapreduce.Job: Running job: job_1516707793508_0012
18/01/23 19:16:56 INFO mapreduce.Job: Job job_1516707793508_0012 running in uber mode : false
18/01/23 19:16:56 INFO mapreduce.Job: map 0% reduce 0%
18/01/23 19:17:02 INFO mapreduce.Job: map 100% reduce 0%
18/01/23 19:17:10 INFO mapreduce.Job: map 100% reduce 100%
18/01/23 19:17:10 INFO mapreduce.Job: Job job_1516707793508_0012 completed successfully
18/01/23 19:17:10 INFO mapreduce.Job: Counters: 49

  File System Counters
    FILE: Number of bytes read=39
    FILE: Number of bytes written=404383
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=162
    HDFS: Number of bytes written=18
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2

  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4134
    Total time spent by all reduces in occupied slots (ms)=4761
    Total time spent by all map tasks (ms)=4134
    Total time spent by all reduce tasks (ms)=4761
    Total vcore-milliseconds taken by all map tasks=4134
    Total vcore-milliseconds taken by all reduce tasks=4761
    Total megabyte-milliseconds taken by all map tasks=4233216
    Total megabyte-milliseconds taken by all reduce tasks=4875264

  Map-Reduce Framework
    Map input records=1
    Map output records=16
    Map output bytes=232
    Map output materialized bytes=39
    Input split bytes=112
    Combine input records=16
    Combine output records=2
    Reduce input groups=2
    Reduce shuffle bytes=39
    Reduce input records=2
    Reduce output records=2
    Spilled Records=4
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=143
    CPU time spent (ms)=1210
    Physical memory (bytes) snapshot=346324992
    Virtual memory (bytes) snapshot=4127256576
    Total committed heap usage (bytes)=222429184

  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

  File Input Format Counters
    Bytes Read=50
  File Output Format Counters
    Bytes Written=18

The mean is: 5.25
[acadgild@localhost mapreduce]$
```

wordmedian: A map/reduce program that counts the median length of the words in the input files.

```

[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordmedian /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs/out/wordmedianfile1
18/01/23 19:19:25 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/01/23 19:19:26 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/01/23 19:19:27 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with
coolRunner to remedy this.
18/01/23 19:19:27 INFO input.FileInputFormat: Total input files to process : 1
18/01/23 19:19:27 INFO mapreduce.JobSubmitter: number of splits:1
18/01/23 19:19:27 INFO Configuration.deprecation: yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enab
d
18/01/23 19:19:28 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1516707793508_0013
18/01/23 19:19:28 INFO impl.YarnClientImpl: Submitted application application_1516707793508_0013
18/01/23 19:19:28 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1516707793508_0013/
18/01/23 19:19:28 INFO mapreduce.Job: Running job: job_1516707793508_0013
18/01/23 19:19:38 INFO mapreduce.Job: Job job_1516707793508_0013 running in uber mode : false
18/01/23 19:19:38 INFO mapreduce.Job: map 0% reduce 0%
18/01/23 19:19:44 INFO mapreduce.Job: map 100% reduce 0%
18/01/23 19:19:52 INFO mapreduce.Job: map 100% reduce 100%
18/01/23 19:19:52 INFO mapreduce.Job: Job job_1516707793508_0013 completed successfully
18/01/23 19:19:52 INFO mapreduce.Job: Counters: 49

File System Counters
  FILE: Number of bytes read=56
  FILE: Number of bytes written=404117
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=162
  HDFS: Number of bytes written=20
  HDFS: Number of read operations=6
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=4008
  Total time spent by all reduces in occupied slots (ms)=4427
  Total time spent by all map tasks (ms)=4008
  Total time spent by all reduce tasks (ms)=4427
  Total vcore-milliseonds taken by all map tasks=4008
  Total vcore-milliseonds taken by all reduce tasks=4427
  Total megabyte-milliseonds taken by all map tasks=4104192
  Total megabyte-milliseonds taken by all reduce tasks=4533248

Map-Reduce Framework
  Map input records=1
  Map output records=8
  Map output bytes=64
  Map output materialized bytes=56
  Input split bytes=112
  Combine input records=8
  Combine output records=5
  Reduce input groups=5
  Reduce shuffle bytes=56
  Reduce input records=5
  Reduce output records=5
  Spilled Records=10
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=129
  CPU time spent (ms)=1170
  Physical memory (bytes) snapshot=346796032
  Virtual memory (bytes) snapshot=4127256576
  Total committed heap usage (bytes)=222429184

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=50
File Output Format Counters
  Bytes Written=20
The median is: 5
[acadgild@localhost mapreduce]$

```

wordstandarddeviation: A map/reduce program that counts the standard deviation of the length of the words in the input files.

```
[acadm@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.9.0.jar wordstandarddeviation /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs/out/wordstandarddeviationfile1
18/01/23 19:22:22 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/01/23 19:22:23 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/01/23 19:22:24 INFO input.FileInputFormat: Total input files to process : 1
18/01/23 19:22:24 WARN hdfs.DataStreamer: Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at org.apache.hadoop.hdfs.DataStreamer.closeResponder(DataStreamer.java:980)
    at org.apache.hadoop.hdfs.DataStreamer.endBlock(DataStreamer.java:630)
    at org.apache.hadoop.hdfs.DataStreamer.run(DataStreamer.java:807)
18/01/23 19:22:24 INFO mapreduce.JobSubmitter: number of splits:1
18/01/23 19:22:25 INFO Configuration.deprecation: yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
18/01/23 19:22:25 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1516707793508_0014
18/01/23 19:22:25 INFO impl.YarnClientImpl: Submitted application application_1516707793508_0014
18/01/23 19:22:25 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1516707793508_0014/
18/01/23 19:22:25 INFO mapreduce.Job: Running job: job_1516707793508_0014
18/01/23 19:22:35 INFO mapreduce.Job: Job job_1516707793508_0014 running in uber mode : false
18/01/23 19:22:35 INFO mapreduce.Job: map 0% reduce 0%
18/01/23 19:22:42 INFO mapreduce.Job: map 100% reduce 0%
18/01/23 19:22:49 INFO mapreduce.Job: map 100% reduce 100%
18/01/23 19:22:50 INFO mapreduce.Job: Job job_1516707793508_0014 completed successfully
18/01/23 19:22:50 INFO mapreduce.Job: Counters: 49

  File System Counters
    FILE: Number of bytes read=56
    FILE: Number of bytes written=404603
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=162
    HDFS: Number of bytes written=29
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2

  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=4086
    Total time spent by all reduces in occupied slots (ms)=4728
```

```
    Total time spent by all map tasks (ms)=4086
    Total time spent by all reduce tasks (ms)=4728
    Total vcore-milliseconds taken by all map tasks=4086
    Total vcore-milliseconds taken by all reduce tasks=4728
    Total megabyte-milliseconds taken by all map tasks=4184064
    Total megabyte-milliseconds taken by all reduce tasks=4841472

  Map-Reduce Framework
    Map input records=1
    Map output records=24
    Map output bytes=352
    Map output materialized bytes=56
    Input split bytes=112
    Combine input records=24
    Combine output records=3
    Reduce input groups=3
    Reduce shuffle bytes=56
    Reduce input records=3
    Reduce output records=3
    Spilled Records=6
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=139
    CPU time spent (ms)=1330
    Physical memory (bytes) snapshot=345300992
    Virtual memory (bytes) snapshot=4127256576
    Total committed heap usage (bytes)=222429184

  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

  File Input Format Counters
    Bytes Read=50
  File Output Format Counters
    Bytes Written=29
The standard deviation is: 2.436698586202241
[acadm@localhost mapreduce]$
```